

clarity for performance. Although this is a common trade-off, it's one I'd r

In his “Simple Made Easy” keynote at the Strange Loop conference, R creator of Clojure, reintroduced an arcane word, *complect*: to join by wea together; to interweave. Imperative programming often forces you to

---

Shifti

tasks so that you can fit them all within a single loop, for efficiency. I programming via higher-order functions such as `map()` and `filter()` allow your level of abstraction, seeing problems with better clarity. I show ma functional thinking as a powerful antidote to incidental complecting.

## Aligning with Language Trends

Michael Feathers, author of *Working with Legacy Code*, captured a key tween functional and object-oriented abstractions in 140 lowly charact

OO makes code understandable by encapsulating moving parts. FP makes standable by minimizing moving parts.

— Mich

Think about the things you know about object-oriented programmin

“moving parts.” Rather than build mechanisms to *control* mutable state, functional languages try to *remove* mutable state, a “moving part.” The theory follows that if a language exposes fewer potentially error-prone features, it is less likely to make errors. I will show numerous examples throughout of functional programming eliminating variables, abstractions, and other moving parts.

In object-oriented imperative programming languages, the units of reuse are the classes and the messages they communicate with, captured in a class diagram. In functional programming, work in that space, *Design Patterns: Elements of Reusable Object-Oriented Software* (by Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides), includes a class diagram with each pattern. In the OOP world, developers are encouraged to build unique data structures, with specific operations attached in the form of methods. In functional programming languages don't try to achieve reuse in quite the same way. In functional programming languages, the preference is for a few key data structures (such as `list`, `set`, and `map`) with highly optimized operations on those data structures. When developers utilize this machinery, developers pass data structures plus higher-order functions that plug into the machinery, customizing it for a particular use.

Consider a simplified portion of the code from **Example 1-2**:

```
regexToList(words, "\\b\\w+\\b").stream()
```

tactically sugared (`w → !NON_WORDS.contains(w)`). The machinery achieves the filtering criteria in an efficient way, returning the filtered list.

Encapsulation at the function level allows reuse at a more granular, functional level than building new class structures for every problem. Dozens of XML libraries exist in the Java world, each with its own internal data structures. One advantage of functional higher-level abstractions is already appearing in the Clojure space. Recent variations in Clojure's libraries have managed to rewrite the `map` function

ically parallelizable, meaning that all `map` operations get a performance boost without developer intervention.

Functional programmers prefer a few core data structures, building operations on top of them to understand them. Object-oriented programmers tend to create many data structures and attendant operations constantly—building new classes and methods for each problem. Encapsulating all logic within classes discourages reuse at the method level, preferring larger frameworks for reuse. Functional programming constructs make it easier to reuse code at the method level.

Consider the `indexOfAny()` method, from the popular Java framework **Guava**, which provides a slew of helpers for Java, in **Example 1-3**.



```

    Some(result.head)
  }
}

```

In **Example 1-5**, I create an indexed version of the input string. Scala's `zip` method takes a collection (the range of numbers up to the length of my input string) and combines it with the collection of `String` characters, creating a new collection of pairs from the original collection. For example, if my input string is `zabcdefghijklmnopqrstuvwxyz`, `zipInput` contains `Vector ((0,z), (1,a), (2,b), (3,y), (4,c), (5,v), (6,n), (7,x))`. The `zip` name comes from the result; it looks as if the two collections are joined like the teeth of a zipper.

Once I have the indexed collection, I use Scala's `for()` comprehension to iterate over the collection of search characters, then access each pair from the indexed collection. Scala allows shorthand access to collection elements, so I can compare the first character to the second item for the collection (`if (char == pair._2)`). If the characters match, I return the index portion of the pair (`pair._1`).

A common source of confusion in Java is the presence of `null`: is it a `String` value or does it represent the absence of a value? In many functional languages (not included), that ambiguity is avoided by the `Option` class, which contains either `None` indicating no return, or `Some`, containing the returned values. For **Example 1-6**, `find` asks for only the first match, so I return `result.head`, the first element of the results collection.

Learning a new programming language is easy: you merely learn the new language's familiar concepts. For example, if you decide to learn JavaScript, your first step is to find a resource that explains how JavaScript's `if` statement works. Typically, developers learn new languages by applying what they know about existing languages. Learning a new paradigm is difficult—you must learn to see different solutions to the same problems.

Writing functional code doesn't require a shift to a functional programming language such as Scala or Clojure but rather a shift in the way you approach problems.

## A Common Example

Once garbage collection became mainstream, it simultaneously eliminated categories of hard-to-debug problems and allowed the runtime to manage memory in a way that is complex and error-prone for developers. Functional programming does the same thing for the algorithms you write, allowing you to work at a higher level of abstraction while freeing the runtime to perform sophisticated optimizations. You can receive the same benefits of lower complexity and higher performance that garbage collection provides, but at a more intimate level, in the way you devise algorithms.

## Imperative Processing

*Imperative programming describes a style of programming modeled after a sequence of commands (imperatives) that modify state. A traditional `for` loop is a good example of the imperative style of programming: establish an initial state, then execute a series of commands for each iteration of the loop.*

To illustrate the difference between imperative and functional programming, we'll start with a common problem and its imperative solution. Let's say that you

```
for(int i = 0; i < listOfNames.size(); i++) {  
    if (listOfNames.get(i).length() > 1) {  
        result.append(capitalizeString(listOfNames.get(i))).append(", ");  
    }  
}  
return result.substring(0, result.length() - 1).toString();  
}
```

result, along with a trailing comma. The last name in the final string strips the comma, so I strip it off the final return value.

Imperative programming encourages developers to perform operations. In this case, I do three things: *filter* the list to eliminate single characters, *transform* the list to capitalize each name, then *convert* the list into a single string. For these three operations *Useful Things* to do to a list. In an imperative language, I use the same low-level mechanism (iteration over the list) for all three operations. Functional languages offer specific helpers for these operations.

## Functional Processing

Functional programming describes programs as expressions and tries to model mathematical formulas, and tries to avoid mutable state. Functional programming languages categorize problems differently than imperative languages. The logical categories listed earlier (*filter*, *transform*, and *convert*) are representations that implement the low-level transformation but rely on the developer to customize the low-level machinery with a higher-order function, supplying parameters. Thus, I could conceptualize the problem as the pseudocode

```
-> filter(x < length < 4)
-> transform(x.capitalize)
-> convert(x + "," + y)
```

Functional languages allow you to model this conceptual solution without worrying about the details.

Consider the company process example from [Example 2-1](#) implemented in [Example 2-3](#).



```
val employees = List("neal", "s", "stu", "j", "rich", "bob", "aiden",  
                    "liam", "mason", "noah", "lucas", "jacob", "jayden", "jack")
```

```
val result = employees  
    .filter(_ .length() > 1)  
    .map(_ .capitalize)  
    .reduce(_ + ", " + _)
```

The use of Scala in [Example 2-3](#) reads much like the pseudocode in [Example 2-2](#), with only the necessary implementation details. Given the list of names, I first filter out the single characters. The output of that operation is then fed into the `map` function, which executes the supplied code block on each element of the collection, returning a new transformed collection. Finally, the output collection from `map` flows to the `reduce` function, which combines each element based on the rules supplied in the lambda expression. In this case, to combine the first two elements, I concatenate them with a comma and space. Using three of these small functions, I don't care what the parameters are named, so [Scala allows me to skip the names and use an underscore instead](#). In the case of `reduce`, I still pass two parameters, which is the expected signature even though I use the same generic indicator, the underscore.

I chose Scala as the first language to show this implementation because of its familiar syntax and the fact that Scala uses industry-consistent names for its collection types. In fact, Java 8 has these same features, and closely resembles the Scala version.

*Example 2-4. Java 8 version of the Company process*

```
public String cleanNames(List<String> names) {  
    if (names == null) return "";  
    return names  
        .stream()  
        .filter(name -> name.length() > 1)  
        .map(name -> capitalize(name))  
        .collect(Collectors.joining(", "));  
}
```

in [Example 2-4](#), I use the `collect()` method rather than `reduce()` because it is more efficient with the Java `String` class; `collect()` is a special case for `reduce()`. Otherwise, it reads remarkably similarly to the Scala code in [Example 2-4](#).

If I was concerned that some of the items in the list might be null, I could add another criterion to the stream:

```
return names
    .stream()
    .filter(name -> name != null)
    .filter(name -> name.length() > 1)
    .map(name -> capitalize(name))
    .collect(Collectors.joining(","));
```

The Java runtime is intelligent enough to combine the null check and length check into a single operation, allowing you to express the idea succinctly yet still have the same code.

Groovy has these features but names them more consistently than scripts such as Ruby. The Groovy version of the “company process” from [Example 2-4](#) is in [Example 2-5](#).

#### *Example 2-5. Processing in Groovy*

```
public static String cleanUpNames(listOfNames) {
    listOfNames
        .findAll { it.length() > 1 }
        .collect { it.capitalize() }
        .join ','
}
```

While [Example 2-5](#) is structurally similar to the Scala code in [Example 2-4](#), the names and substitution identifier differ. Groovy’s `findAll` on a collection



tions and stop going immediately for detailed implementations.

What are the benefits of thinking at a higher level of abstraction? First, you to categorize problems differently, seeing commonalities. Second, it's time to be more intelligent about optimizations. In some cases, reordering a stream makes it more efficient (for example, processing fewer items) if it achieves the ultimate outcome. Third, it allows solutions that aren't possible when you're elbow deep in the details of the engine. For example, consider the code required to make the Java code in [Example 2-1](#) run across multiple threads. If you control the low-level details of iteration, you must weave the thread code into the logic. In the Scala version, I can make the code parallel by adding `par` to the stream as in [Example 2-8](#).

*Example 2-8. Scala processing in parallel*

```
val parallelResult = employees
    .par
    .filter(_ .length() > 1)
    .map(_ .capitalize)
    .reduce(_ + "," + _)

public String cleanNamesP(List<String> names) {
    if (names == null) return "";
    return names
        .parallelStream()
        .filter(n -> n.length() > 1)
        .map(e -> capitalize(e))
        .collect(Collectors.joining(","));
}
```

required to write the `sum()` method—in Java 8, it is one of the stream methods that generates values.

In physics, energy is differentiated into *potential*, energy stored and ready to be used, and *kinetic*, energy expenditure. For collections in languages such as Java before Java 8, all collections acted as kinetic energy: the collection resolved values immediately, entering no intermediate state. Streams in functional languages work more like potential energy, which is stored for later use. The stream holds the origin (in [Example 2-12](#), the origin comes from the `range()` method) and whatever operations have been attached to the stream, such as filtering operations. The stream transitions from potential to kinetic until the developer “asks” for values, using a terminal operation such as `forEach()` or `sum()`. The stream is passable as a parameter to other methods, and additional criteria added later to its potential until it becomes kinetic. This is the concept of [lazy evaluation](#), which I cover in detail in [Chapter 4](#).

This style of coding was possible, with difficulty, in previous versions of Java, but not in useful frameworks.

## Functional Java Number Classifier

While all modern languages now include higher-order functions, many developers stay on older versions of runtimes such as Java for many years for no good reason. [Functional Java](#) is an open source framework whose mission includes adding many functional idioms to Java post version 1.5 with as little intrusiveness as possible. For example, because the Java 1.5-era JDK don't include higher-order functions, functional Java mimics their use via generics and anonymous inner classes. A number classifier takes on a different look when implemented using Functional Java, as shown in [Example 2-13](#).

*Example 2-13. Number classifier using the Functional Java framework*

In [Example 2-13](#), I used the `foldLeft()` method, which collapses elements toward the first element. For addition, which is commutative, the direction doesn't matter. If, on the other hand, I need to use an operation in which order is important, I also use a `foldRight()` variant.



Higher-order abstractions eliminate friction.

When we were talking about the waning of Smalltalk versus the waxing of Java, I did some extensive work in both and says that he initially viewed the switch from Smalltalk to Java as a syntactic inconvenience, but eventually as an impediment to the kind of thinking afforded in the previous world. Placing syntactic hurdles around the use of higher-order abstractions adds needless friction to the thought process.



Don't add needless friction.

## Filter

A common operation on lists is filtering: creating a smaller list by filtering out elements from a list based on some user-defined criteria. Filtering is illustrated in [Figure 2-14](#).



situation keyword `cc` as a placeholder, and the last line of the method returns the value, which is the list of factors in this case.



Use `filter` to produce a subset of a collection based on some filtering criteria.

## Map

The `map` operation transforms a collection into a new collection by applying a function to each of the elements, as illustrated in **Figure 2-2**.

---

Common Bu

`aliquot-sum` in each case and returns the appropriate keyword (an element is lineated with a leading colon).



Use `map` to transform collections in situ.

## Fold/Reduce

The third common function has the most variations in name, and many differences, among popular languages. `foldLeft` and `reduce` are specific variations of a manipulation concept called *catamorphism*, which is a generalization of

The `reduce` and `fold` operations have overlapping functionality, with subtle differences from language to language. Both use an accumulator to gather values. The `reduce` function is generally used when you need to supply an initial value for the accumulator, whereas `fold` starts with nothing in the accumulator. The order of operation on the collection is specified by the specific method name (for example, `foldLeft` and `foldRight`). Neither of these operations mutates the collection.

I show the `foldLeft()` function in Functional Java. In this case, a “fold” operation

```
1 // Example 2-19: A fold operation to calculate the aliquot sum
2 // The aliquot sum of a number is the sum of its proper divisors.
3 // For example, the aliquot sum of 12 is 1 + 2 + 3 + 4 + 6 = 16.
4 // The aliquot sum of 1 is 0.
5 // The aliquot sum of a prime number is 1.
6 // The aliquot sum of 0 is 0.
7 // The aliquot sum of a negative number is 0.
8 // The aliquot sum of a non-integer is 0.
9 // The aliquot sum of a non-numeric value is 0.
10 // The aliquot sum of a null value is 0.
```

At first it's not obvious how the one-line body in **Example 2-19** performs a fold operation to calculate the `aliquotSum`. In this case, the *fold* operation refers to a *fold left* operation that combines each element of the list with the next one, accumulating the result for the entire list. A fold left combines the list elements leftward, starting with an initial value and accumulating each element of the list in turn to yield a final result. **Example 2-19** illustrates a fold operation.

```
1 // Example 2-19: A fold operation to calculate the aliquot sum
2 // The aliquot sum of a number is the sum of its proper divisors.
3 // For example, the aliquot sum of 12 is 1 + 2 + 3 + 4 + 6 = 16.
4 // The aliquot sum of 1 is 0.
5 // The aliquot sum of a prime number is 1.
6 // The aliquot sum of 0 is 0.
7 // The aliquot sum of a negative number is 0.
8 // The aliquot sum of a non-integer is 0.
9 // The aliquot sum of a non-numeric value is 0.
10 // The aliquot sum of a null value is 0.
```

Because addition is commutative, it doesn't matter if you do a `foldLeft()` or `foldRight()`. But some operations (including subtraction and division) cannot be commutative, so the `foldRight()` method exists to handle those cases. In purely functional programming, left and right folds have implementation differences. For example, right folds can operate on infinite lists whereas left folds cannot.

**Example 2-13** uses the Functional Java-supplied `add` enumeration; the `add` enumeration includes the most common mathematical operations for you. But what

to produce a different-sized (usually smaller but not necessarily) value, a collection or a single value.



Use `reduce` or `fold` for piecewise collection processing.

`collect()`, `map()`, and `inject()` appear. Once you become accustomed to these tools in your toolbox, you'll find yourself turning to them again and again.

One of the challenges of learning a different paradigm such as functional programming is learning the new building blocks and “seeing” them peek out of problem solution. In functional programming, you have far fewer abstractions, more generic (with specificity added via higher-order functions). Because functional programming relies heavily on passed parameters and composition, you have to learn about the interactions among moving parts, making your job easier.

## Synonym Suffering



This second version is less verbose because Scala allows substitution of `p` underscores. Both versions yield the same result.

Many examples of filtering operations use numbers, but `filter()` applies to **any collection**. This example applies `filter()` to a list of words to determine which words are three letters long:

```
val words = List("the", "quick", "brown", "fox", "jumped",  
                 "over", "the", "lazy", "dog")  
words filter (_.length == 3)  
// List(the, fox, the, dog)
```

Another filtering variant in Scala is the `partition()` function, which returns a **partitioned version** of a collection by splitting it into multiple parts; the original collection is **unchanged**. The split is based on the higher-order function that you pass as the separation criteria. Here, the `partition()` function returns two lists according to which list members are divisible by 3:

```
numbers partition (_ % 3 == 0)  
// (List(3, 6, 9), List(1, 2, 4, 5, 7, 8, 10))
```

The `filter()` function returns a collection of matching elements, while `find()` returns **only the first match**:

```
numbers find (_ % 3 == 0)  
// Some(3)
```

However, the **return value for `find()`** isn't the matched value itself, but is **wrapped in an `Option` class**. `Option` has two possible values: `Some` or `None`. In some other functional languages, `Option` is used as a convention to avoid `null` in the absence of a value. The `Some()` instance wraps the actual return value, 3 in the case of `numbers find (_ % 3 == 0)`. If I try to find something that doesn't exist, the return is `None`:

I discuss `Option` and similar classes in depth in [Chapter 5](#).

Scala also includes several functions that process a collection based on a predicate and either retain values or discard them. The `takeWhile()` function returns the largest set of values from the front of the collection that satisfy the predicate:

```
List(1, 2, 3, -4, 5, 6, 7, 8, 9, 10) takeWhile (_ > 0)  
// List(1, 2, 3)
```

The `dropWhile()` function skips the largest number of elements that satisfy the predicate:

```
words dropWhile (_ startsWith "t")  
// List(quick, brown, fox, jumped, over, the, lazy, dog)
```

## Map

The second major functional transformation that's common to all functional languages is `map`. A map function accepts a higher-order function and a collection, and it passes the function to each element and returns a collection of the results. The return value (unlike with filtering) is the same size as the original collection, but with

## Scala

Scala's `map()` function accepts a code block and returns the transformed

```
List(1, 2, 3, 4, 5) map (_ + 1)
// List(2, 3, 4, 5, 6)
```

The `map()` function works on all applicable types, but it doesn't necessarily return a transformed version of each element of the collection. In this example, it returns the sizes of all elements in a string:

```
words map (_.length)
// List(3, 5, 5, 3, 6, 4, 3, 4, 3)
```

Nested lists occur so often in functional programming languages that a function for denesting—typically called *flattening*—is common. Here is an example of a nested list:

```
List(List(1, 2, 3), List(4, 5, 6), List(7, 8, 9)) flatMap (_.toList)
```

## Scala

Scala has the richest set of fold operations, in part because it facilitates scenarios that don't appear in the dynamically typed Groovy and Clojure. Here are two commonly used to perform sums:

The `reduceLeft()` function assumes that the first element is the left side of the operation. For operators such as plus, the placement of operands is irrelevant. For operators such as divide, the placement of operands matters. If you want to reverse the order of the operator, use `reduceRight()`:

```
List.range(1, 10) reduceRight(_ - _)
// 0 - 1 = -1
```



Thus, it applies  $8 - 9$  first, then uses that result as the *second* parameter for subsequent calculations.

Understanding when you can use higher-level abstractions such as `reduceLeft` is a key to mastery of functional programming. This example uses `reduceLeft` to determine the longest word in a collection:

```
words.reduceLeft((a, b) => if (a.length > b.length) a else b)
// jumped
```

The `reduce` and `fold` operations have overlapping functionality, with subtle differences that I discussed earlier. One obvious difference suggests common use cases. The signature of `reduceLeft[B >: A](op: (B, A) => B): B` shows that the type of the result that's expected is the function to combine elements. The initial value is the first value in the collection. In contrast, the signature of `foldLeft[B, A](initial: B, op: (B, A) => B): B` indicates an initial seed value for the result, so you can use a value that is different from the type of the list elements.

Here's an example of summing a collection by using `foldLeft()`:

```
List.range(1, 10).foldLeft(0)(_ + _)
// 45
```

Scala supports operator overloading, so the two common fold operations, `foldRight` and `foldLeft`, have corresponding operators: `/:` and `\:`, respectively. Thus, you can write a terser version of `sum` by using `foldLeft`:

```
(0 /: List.range(1, 10)) (_ + _)
// 45
```

Many of the features of functional languages were commonplace a decade ago but make perfect sense now because they optimize developer productivity.

One of the values of functional thinking is the ability to cede control of low-level details (such as garbage collection) to the runtime, eliminating a swath of bugs and complexity. While many developers are accustomed to blissful ignorance for bedrock concerns such as memory, they are less accustomed to similar abstractions appearing at a higher level. Yet these higher-level abstractions serve the same purpose: handling

In this chapter, I show five ways developers in functional languages can surrender control to the language or runtime, freeing themselves to work on more relevant problems.

## Iteration to Higher-Order Functions

I already illustrated the first example of surrendering control in [Example 1.1](#). Iteration with functions such as `map`. What's ceded here is clear: if you can perform an operation you want to perform in a higher-order function, the language does it for you efficiently, including the ability to parallelize the operation with the added power of a parallel modifier.

high performance, even with the new `Stream` API in Java 8. Once you understand the power of iteration, however, you can apply that power in a more succinct way.



Always understand one level below your normal abstraction.

Programmers rely on abstraction layers to be effective: no one program manipulates bits of memory, but delivers values just 0 and 1. Abstraction

## Closures

All functional languages include closures, yet this language feature is often almost mystical terms. A closure is a function that carries an implicit binding of variables referenced within it. In other words, the function (or method) text around the things it references.

Here is a simple example, written in Groovy, of creating a closure block shown in [Example 3-1](#).

*Example 3-1. Simple closure binding in Groovy*

```
class Employee {  
    def name, salary  
}  
  
def paidMore(amount) {  
    return {Employee e -> e.salary > amount}  
}  
  
isHighPaid = paidMore(100000)
```



In Example 3-1, I define a simple `Employee` class with two fields. The `paidMore` function that accepts a parameter amount. The return of this function is a code block, or *closure*, accepting an `Employee` instance. The type declaration serves as useful documentation in this case. I can assign this code block to `isHighPaid`, supplying the parameter value of `100,000`. When I make the assignment, I bind the value of `100,000` to this code block forever. Thus, when I evaluate via this code block, it will opine about their salaries by using the permanent value, as shown in Example 3-2.

*Example 3-2. Executing the closure block*

```
def Smithers = new Employee(name:"Fred", salary:120000)
def Homer = new Employee(name:"Homer", salary:80000)
println isHighPaid(Smithers)
println isHighPaid(Homer)
// true, false
```

In Example 3-2, I create a couple of employees and determine if their salaries are high by a certain criterion. When a closure is created, it creates an enclosure around the code block that is enclosed within the scope of the code block (thus the name *closure*). Each closure block has unique values, even for private variables. For example, I can create another instance of my `paidMore` closure with another binding (and the same assignment), as shown in Example 3-3.

*Example 3-3. Another closure binding*

```
isHigherPaid = paidMore(200000)
println isHigherPaid(Smithers)
println isHigherPaid(Homer)
def Burns = new Employee(name:"Monty", salary:1800000)
println isHigherPaid(Burns)
// false, false, true
```

Closures are used quite often as a portable execution mechanism in functions and frameworks, passed to higher-order functions such as `map()` as the code. Functional Java uses anonymous inner classes to mimic some of the behavior, but they can't go all the way because Java before version 8 did not have closures. But what does that mean?

Example 3-4 shows an example of what makes closures so special.

**Example 3-4.** When the closure block is garbage collected, all the references it contains are reclaimed as well.

It's a bad idea to create a closure just so that you can manipulate its internal state. This is done here to illustrate the inner workings of closure bindings. Binding to immutable values (as shown in **Example 3-1**) is more common.

The closest you could come to the same behavior in Java prior to Java 8 is a language that has functions but not closures, appears in **Example 3-5**.

Several variants of the Counter class are possible (creating anonymous classes, using generics, etc.), but you're still stuck with managing the state yourself. This is why the use of closures exemplifies functional thinking: allow the runtime to manage the state. Rather than forcing yourself to handle field creation and babying the state, the horrifying prospect of using your code in a multithreaded environment is avoided by letting the language or framework invisibly manage that state for you.



Let the language manage state.

Closures are also an excellent example of *deferred execution*. By binding a closure to a block, you can wait until later to execute the block. This turns out to be useful in many scenarios. For example, the correct variables or functions might not be available at definition time but are at execution time. By wrapping the execution context in a closure, you can wait until the proper time to execute it.

Imperative languages use *state* to model programming, exemplified by programming. Closures allow us to model *behavior* by encapsulating both code and a single construct, the closure, that can be passed around like traditional and executed at exactly the correct time and place.



Capture the *context*, not the *state*.

## Currying and Partial Application

Currying and partial application are language techniques derived from (based on work by twentieth-century mathematician Haskell Curry and techniques are present in various types of languages and are omnipresent in languages in one form or another. Both currying and partial application ability to manipulate the number of arguments to functions or methods supplying one or more default values for some arguments (known as fixing). Most functional languages include currying and partial application, but use them in different ways.



to the casual observer, currying and partial application appear to be the same effect. With both, you can create a version of a function with presupplied values for some of the arguments:

- *Currying* describes the conversion of a multiargument function into a chain of single-argument functions. It describes the transformation process of the converted function. The caller can decide how many arguments to supply, thereby creating a derived function with that smaller number of arguments.
- *Partial application* describes the conversion of a multiargument function into a function that accepts fewer arguments, with values for the elided arguments supplied in advance. The technique's name is apt: it partially applies some arguments to a function, returning a function with a signature that consists of the remaining arguments.

With both currying and partial application, you supply argument values in advance, creating a function that's invocable with the missing arguments. But currying a function creates a new function that returns the next function in the chain, whereas partial application binds argument values that you supply during the operation, producing a function with a specific number of arguments. This distinction becomes clearer when you compare currying with arity greater than two.



For example, the fully curried version of the `process(x, y, z)` function is `process(x)(y)(z)`, where both `process(x)` and `process(x)(y)` are functions that accept a single argument. If you curry only the first argument, the return value of `process(x)` is a function that accepts a single argument that in turn accepts a single argument. In contrast, with partial application, you are left with a function of smaller arity. A partial application for a single argument on `process(x, y, z)` yields a function that accepts two arguments: `process(y, z)`.

The two techniques' results are often the same, but the distinction is often misconstrued. To complicate matters further, Groovy implements both partial application and currying but calls both currying. And Scala has both `Function` and the `PartialFunction` class, which are distinct concepts with similar names.

## In Groovy

## Scala

Scala supports currying and partial application, along with a trait that provides the ability to define constrained functions.

### Currying

In Scala, functions can define multiple argument lists as sets of parentheses. If you call a function with fewer than its defined number of arguments, the return value is a function that takes the missing argument lists as its arguments. Consider the example in the Scala documentation that appears in [Example 3-10](#).

*Example 3-10. Scala's currying of arguments*

In [Example 3-11](#), I first create a `price()` function that returns a map of product and price. Then I create a `withTax()` function that accepts a state as an argument. However, within a particular source file, I know that I will work with one state's taxes. Rather than "carry" the extra argument for every call, I can partially apply the state argument and return a version of the function in which the state value is fixed. The `locallyTaxed()` function accepts a single argument

### Partial (constrained) functions

The `Scala PartialFunction` trait is designed to work seamlessly with `Partial`, which is covered in detail in [Chapter 6](#). Despite the similarity in name, `PartialFunction` does not create a partially applied function. Instead, you define a function that works only for a defined subset of values and types.

Case blocks are one way to apply partial functions. [Example 3-12](#) uses `case` without the traditional corresponding `match` operator.

*Example 3-12. Using `case` without `match`*

In [Example 3-12](#), I create a map of city and state correspondence. The `map()` function on the collection, and `map()` in turn pulls apart the key-value pairs and prints them. In Scala, a code block that contains `case` statements is one way to define an anonymous function. You can define anonymous functions more conventionally using `case`, but the `case` syntax provides the additional benefits that [Example 3-12](#) illustrates.

`MatchError` as the function tries to increment the "seven" string. But `collect()` works correctly. Why the disparity and where did the error go?

Case blocks define partial functions, but not partially applied functions. Partial functions have a limited range of allowable values. For example, the mathematical function  $1/x$  is invalid if  $x = 0$ .

Partial functions offer a way to define constraints for allowable values. In the `collect()` invocation in [Example 3-13](#), the case is defined for `Int`, but not for `String`. The "seven" string isn't collected.

In [Example 3-16](#), I define a partial function to accept any type of input (A partial function can react to a subset of types). However, notice that I can also call the `collect()` function for the partial function. Implementers of the `PartialFunction` trait can call `isDefinedAt()`, which is implicitly defined. [Example 3-13](#) shows how `map()` and `collect()` behave differently. The behavior of partial functions is discussed in [Chapter 10](#).

**Partial functions and `collect()`**  
The difference: `collect()` is designed to accept partial functions and to call `isDefinedAt()` function for elements, ignoring those that don't match.

Partial functions and partially applied functions in Scala are similar in that they offer a different set of orthogonal features. For example, nothing prevents you from partially applying a partial function.

## Common Uses



application as well as a piece in test-driven programming.

### Function factories

Currying (and partial application) work well for places where you implement a function in traditional object-oriented languages. To illustrate, [Example 3-17](#) implements a simple adder function in Groovy.

*Example 3-17. Adder and incrementer in Groovy*

[In Example 3-17](#), I use the `add(x, y)` function to create the incrementer.

### Template Method design pattern

One of the Gang of Four design patterns is the Template Method pattern, which helps you to help you define algorithmic shells that use internal abstract methods to provide implementation flexibility. Partial application and currying can solve the problem of using partial application to supply known behavior and leaving the other parts free for implementation specifics mimics the implementation of this design pattern.

I show an example of using partial application and other functional techniques to recreate several design patterns (including Template Method) in [Chapter 4](#).

### Implicit values

When you have a series of function calls with similar argument values, you can use currying to supply implicit values. For example, when you interact with the `HttpURLConnection` framework, you must pass the data source as the first argument. By using currying, you can supply the value implicitly, as shown in [Example 3-18](#).




## Recursion

*Recursion*, which (according to Wikipedia) is the “process of repeating similar way,” is another example of ceding details to the runtime, and is associated with functional programming. In reality, it's a computer-science thing over things by calling the same method from itself, reducing the collection, and always carefully ensuring you have an exit condition. Many times, it's not the easy-to-understand code because the core of your problem is the need to do the thing over and over to a diminishing list.

## Seeing Lists Differently

*Figure 3-1. Lists as indexed slots*

Many functional languages have a slightly different perspective on lists, and Groovy shares this perspective. Instead of thinking of a list as indexed slots, it's as a combination of the first element in the list (the head) plus the remainder (the tail), as shown in [Figure 3-2](#).



*Figure 3-2. A list as its head and tail*

Thinking about a list as head and tail allows me to iterate through it using `head` and `tail`, as shown in [Example 3-20](#).

If not, I print out the first element in the list, available via Groovy's `head` property, then recursively call the `recurseList()` method on the remainder of the list.

Recursion often has technical limits built into the platform, so this is not a panacea. But it should be safe for lists that contain a small number of items.

The difference between Examples 3-21 and 3-22 highlights an important distinction. Who's minding the state? In the imperative version, *I* am. *I* must create a new list named `new_list`, *I* must add things to it, and *I* must return it when I'm done. In the recursive version, the *language* manages the return value, building it up from the recursive return for each method invocation. Notice that every call to the `filter()` method in Example 3-22 is a return call, which builds up the return value on the stack. You can cede responsibility for `new_list`; the language will manage it for you.



Recursion allows you to cede state management to the runtime.

The list-construction operators in Scala make the return conditions for `filter()` readable and easy to understand. The code in Example 3-23 is one of the examples in both recursion and currying from the Scala documentation. The `filter()` method recursively filters a list of integers via the parameter `p`, a predicate function. This is a term in the functional world for a Boolean function. The `filter()` method sees if the list is empty and, if it is, simply returns; otherwise, it checks the first element in the list (`xs.head`) via the predicate to see if it should be included in the result.

## Tail-Call Optimization

One of the major reasons that recursion isn't a more commonplace operation is stack growth. Recursion is generally implemented to place intermediate results on the stack, and languages not optimized for recursion will suffer stack overflow. Languages such as Scala and Clojure have worked around this limitation in various ways. One way that developers can help runtimes handle this problem is tail-call optimization. When the recursive call is the last call in the function, runtimes can optimize by returning results on the stack rather than force it to grow.

Many functional languages (such as **Erlang**), implement tail recursion optimization to prevent stack growth. Tail recursion is used to implement long-running Erlang processes.

My guess is that you don't use recursion at all now—it's not even a part of the standard library. However, part of the reason lies in the fact that most imperative languages don't support it, making it more difficult to use than it should be. By adding first-class functions and support, functional languages make recursion a candidate for simple solutions.

## Streams and Work Reordering

before `filter()`. When thinking imperatively, the instinct is to place the filtering operation before the mapping operation, so that the `map` has less work to process a smaller list. However, many functional languages (including Java 8 and the Functional Java framework) define a `Stream` abstraction. A `Stream` acts like a collection, but it has no backing values, and instead uses a stream of values from a source to a destination. In [Example 3-24](#), the source is the `names` collection, and the destination (or “terminal”) is `collect()`. Between these operations, both `filter()` and `map()` are *lazy*, meaning that they defer execution as long as possible. They don’t try to produce results until a downstream terminal “asks” for them.

For the lazy operations, intelligent runtimes can reorder the result of the operations to be more efficient. In [Example 3-24](#), the runtime can flip the order of the lazy operations to be more efficient, performing the filtering before the mapping. As with many other additions to Java, you must ensure that the lambda blocks you pass to `filter()` and `map()` don’t have side effects, which will lead to unpredictable results.

Allowing the runtime to optimize when it can is another great example of giving away mundane details and focusing on the problem domain. The `Stream` interface is a *functional* implementation of the problem domain.

I discuss laziness in more detail in [Chapter 4](#) and Java 8 streams in [Chapter 5](#).



## Memoization

The word **memoization** was coined by Donald Michie, a British artificial intelligence researcher, to **refer to function-level caching for repeating values**. Today, memoization is common in functional programming languages, either as a built-in feature or as a library, but that's relatively easy to implement.

Memoization helps in the following scenario. Suppose you have a computationally intensive function that you **must call repeatedly**. A common solution is to maintain an **internal cache**. Each time you calculate the value for a certain set of parameters, you store that value in the **cache, keyed to the parameter value(s)**. In the future, if the function is invoked with previous parameters, return the value from the cache rather than recalculate it. Function caching is a classic computer science trade-off: it **uses extra memory** (which we frequently have in abundance) **to achieve better performance**.

**Functions must be *pure* for the caching technique to work.** A *pure* function has no side effects: it references no other mutable class fields, doesn't set any, and returns only the return value, and relies only on the parameters for input. All the functions in the `java.lang.Math` class are excellent examples of pure functions. Obviously, you can reuse cached results successfully only if the function reliably returns the same value for a given set of parameters.

## Adding Memoization

Functional programming strives to minimize moving parts by building mechanisms into the runtime. Memoization is a feature built into a program that enables automatic caching of recurring function-return values. In automatically supplies the code I've written in Examples 4-1 and 4-3. Languages support memoization, including Groovy.

In order to memoize a function in Groovy, you define it as a closure, then use the `memoize()` method to return a function whose results will be cached.

Memoizing a function is a metafunction application: doing something itself rather than the function results. Currying, discussed in Chapter 3, is another example of a metafunction technique. Groovy built memoization into its `Function` interface. Other languages implement it differently.

Memoized	956 ms
Memoized (2nd)	19 ms

Memoizing everything slows down the first run but has the fastest subsequent runs in any case—but only for small sets of numbers. As with the imperative case tested in Example 4-3, large number sets impede performance drastically. The memoized version runs out of memory in the 8,000-number case. But for small sets, the memoized approach to be robust, safeguards and careful awareness of the execution environment are required—another example of imperative moving parts. With memoization, the caching occurs at the function level. Look at the memoization results for 10,000 found in Table 4-5.

Table 4-5. Results for range 1–10,000

`memoizeBetween()` Creates a caching variant of the closure with automatic cache size adjustment and upper limits on the cache size

---

In the imperative version, the developer owns the code (and responsibility for correctness). Languages build generic machinery—sometimes with customization knobs of alternate functions or parameters—that you can apply to standard conditions. These conditions are a fundamental language element, so optimizing at that level gives you more functionality for free. The memoization versions in this chapter with small overheads will outperform the handwritten caching code handily. In fact, I'll never be as efficient as the language designers can because they can bend the rules. Language designers have access to low-level parts that developers don't have. Optimization opportunities beyond the grasp of mere mortals. Not only can they handle caching more efficiently, I want to cede that responsibility to them so I can think about problems at a higher level of abstraction.



Language designers will always build more efficient mechanisms because they are allowed to bend rules.

Building a cache by hand is straightforward, but it adds statefulness and complexity to the code. Using functional-language features like memoization, I can achieve the same function level, achieving better results (with virtually no change to my imperative version). Functional programming eliminates moving parts, so you can focus your energy on solving real problems.

Of course, you don't have to rely on an existing class to layer memoization

the memoized function via an explicit invocation to `call()`.

Most functional languages either include memoization or make it trivial. For example, memoization is built into Clojure; you can memoize any function via the built-in (`memoize`) function. For example, if you have an existing (hash) function, you can memoize it via (`memoize (hash "homer")`) for a caching version. Clojure also implements the name-hashing algorithm from **Example 4-6** in Clojure.



call() method. In the Clojure version, the memoized method call is on the surface, with the added indirection and caching invisible to the

Scala doesn't implement memoization directly but has a collection method `getOrElseUpdate()` that handles most of the work of implementing

Example 4-9.

*Example 4-9. Memoization implementation in Scala*

```
def memoize[A, B](f: A => B) = new (A => B) {  
  val cache = scala.collection.mutable.Map[A, B]()  
  def apply(x: A): B = cache.getOrElseUpdate(x, f(x))  
}  
  
def nameHash = memoize(hash)
```

The `getOrElseUpdate()` function in Example 4-9 is the perfect operator for a cache: it either retrieves the matching value or creates a new entry when

It is worth reiterating the importance of immutability for anything you memoize. A memoized function relies on anything other than parameters to generate results, and it will receive unpredictable outcomes. If your memoized function has side effects, it won't be able to rely on that code executing when the cached value is re-



Make sure all memoized functions:

- Have no side effects
- Never rely on outside information

As runtimes become more sophisticated and we have plenty of machinery at our disposal, advanced features such as memoization become common in every mainstream language. For example, although Java 8 doesn't include memoization, it is easy to implement it atop the new lambda features.

## Laziness

Lazy evaluation—deferral of expression evaluation for as long as possible—is a feature of many functional programming languages. Lazy collections deliver their elements as needed rather than precalculating them, offering several benefits. First, they defer expensive calculations until they're absolutely needed. Second, you can work with infinite collections, which keep delivering elements as long as they keep receiving requests. Third, lazy use of functional concepts such as map and filter enable you to write more efficient code. Java doesn't natively support laziness until Java 8, but many frameworks and successor languages do.

Consider the snippet of pseudocode for printing the length of a list in Haskell:

*Example 4-10. Pseudocode illustrating nonstrict evaluation*

```
print length([2+1, 3*2, 1/0, 5-4])
```

If you try to execute this code, the result will vary depending on the type of language it's written in: *strict* or *nonstrict* (also known as *lazy*). In a strict language, executing (or perhaps even compiling) this code results in an exception because of the list's third element. In a nonstrict language, the runtime accurately reports the number of items in the list. After all, the method is called `length()`, not `lengthAndThrowExceptionWhenDivByZero()`! Haskell is a nonstrict language in common use. Alas, Java doesn't support nonstrict evaluation, but you can still take advantage of the concept of laziness in Java by deferring evaluation, and some next-generation languages are lazier than Java by default.

Functional programming languages approach code reuse differently than object-oriented languages. Object-oriented languages tend to have many data structures and many operations, whereas functional languages exhibit few data structures and many operations. Object-oriented languages encourage you to create class-structures, and you can capture recurring patterns for later reuse. Functional languages achieve reuse by encouraging the application of common transformation functions to data structures, with higher-order functions to customize the operation for specific cases.

In this chapter, I cover various ways that languages have *evolved* solutions to recurring problems in software. I discuss the attitudinal change in functional programming.

A seminal work in that space, *Design Patterns: Elements of Reusable Object-Oriented Software* (Addison-Wesley, 1994), includes at least one class diagram with each chapter. In the OOP world, developers are encouraged to create unique data structures and specific operations attached in the form of methods. Functional programming languages don't try to achieve reuse in the same way. They prefer a few key data structures (such as `list`, `set`, and `map`) with highly optimized operations on those data structures. They pass data structures plus higher-order functions to "plug into" the operations.

---

For example, the `filter()` function in several languages accepts a code block as the "plug-in" higher-order function. The function terminates the filter criteria, and the machinery applies the filter criteria to the list, returning the filtered list.

Encapsulation at the function level enables reuse at a more granular, functional level than building custom class structures. One advantage of this approach is its simplicity, appearing in Clojure. For example, consider the case of parsing XML. A number of frameworks exist for this task in Java, each with custom data structures and XML semantics (for example, SAX versus DOM). Clojure parses XML into a simple, uniform structure, and then provides a set of functions to manipulate it.